

FA2 / FA3 ATTENTION VS CUDNN SDPA

NVIDIA RTX PRO 5000 (SM-120A) B=1, H=32, N=16384, D=128

CONFIGURATION	TFLOPS	SPEEDUP	RELATIVE
CUDNN SDPA BASELINE	232.4	1.00X	

■ FA2 MMA

SK=2, F16ACC	157.9	0.68X	
SK=1, F32ACC	145.6	0.63X	
SK=2, F32ACC	166.5	0.72X	

■ FA2 CUTE MMA

SK=1, F32ACC	189.6	0.82X	
SK=2, F32ACC	197.0	0.85X	

■ FA2 TMA WS (1 WG)

SK=1, SV=1, F16ACC	161.2	0.69X	
SK=2, SV=1, F16ACC	189.4	0.81X	
SK=2, SV=2, F16ACC	190.5	0.82X	
SK=2, SV=1, F32ACC	204.9	0.88X	

■ FA3 TMA WS (2 WG)

SK=1, SV=1, F16ACC	210.1	0.90X	
SK=1, SV=1, F32ACC	210.8	0.91X	

■ FA2 CUTE TMA WS (1 WG)

SK=2, SV=1, F32ACC	220.0	0.95X	
SK=3, SV=1, F32ACC	223.4	0.96X	

■ FA2 CUTE TMA WS PERSISTENT-CTA (1 WG)

D=128	242.5	1.04X	
-------	-------	-------	-------------

HGEMM CUTE SWIZZLE VS CUBLAS

NVIDIA RTX PRO 5000 (SM-120A) M=N=K=4096, F16ACC

CONFIGURATION	TFLOPS	SPEEDUP	RELATIVE
CUBLAS BASELINE	236.9	1.00X	

■ HGEMM CUTE SWIZZLE (F16ACC)

S=2, BLK-SW=0	229.5	0.97X	
S=2, BLK-SW=1	231.1	0.98X	
S=3, BLK-SW=0	245.6	1.04X	
S=3, BLK-SW=1	246.7	1.04X	

SPEEDUP IS RELATIVE TO THE VENDOR BASELINE
(CUBLAS FOR HGEMM, CUDNN SDPA FOR ATTENTION).

FA SPLIT-D VS CUDNN SDPA

NVIDIA RTX PRO 5000 (SM-120A) B=1, H=32, N=8192, D=320

CONFIGURATION	TFLOPS	SPEEDUP	RELATIVE
CUDNN SDPA BASELINE	69.8	1.00X	

■ SPLIT-D TMA WS

SK=1, SV=1	86.6	1.24X	
SK=2, SV=2	139.9	2.01X	

■ SPLIT-D TMA NON-WS

SK=2, SV=2	187.9	2.69X	
SK=2, SV=3	188.8	2.71X	
SK=3, SV=2	204.3	2.93X	

SK = K STAGES (FA), SV = V STAGES (FA), S = STAGES (HGEMM)